

Week 5: Graph Algorithms

Directed Graphs

Directed Graphs (Digraphs)


2/86

In our previous discussion of graphs:

- an edge indicates a relationship between two vertices
- an edge indicates nothing more than a relationship

In many real-world applications of graphs:

- edges are directional ($v \rightarrow w \neq w \rightarrow v$)




- edges have a *weight* (cost to go from $v \rightarrow w$)

... Directed Graphs (Digraphs)

3/86

Example digraph and adjacency matrix representation:



	a	b	c	d	e	f	g
a	1	0	0	1	0	0	0
b	1	0	1	0	0	0	0
c	0	1	0	1	0	0	1
d	0	1	0	0	0	1	1
e	0	0	0	1	0	0	0
f	1	0	0	1	0	0	1
g	0	0	0	0	0	0	1

Undirectional \Rightarrow symmetric matrix

Directional \Rightarrow non-symmetric matrix

Maximum #edges in a digraph with V vertices: V^2

... Directed Graphs (Digraphs)

4/86

Terminology for digraphs ...

Directed path: sequence of $n \geq 2$ vertices $v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_n$

- where $(v_i, v_{i+1}) \in \text{edges}(G)$ for all v_i, v_{i+1} in sequence
- if $v_1 = v_n$, we have a *directed cycle*

Reachability: w is reachable from v if \exists directed path v, \dots, w

5/86

Digraph Applications

Potential application areas:

Domain	Vertex	Edge
Web	web page	hyperlink
scheduling	task	precedence
chess	board position	legal move
science	journal article	citation
dynamic data	malloc'd object	pointer
programs	function	function call
make	file	dependency

... Digraph Applications

6/86

Problems to solve on digraphs:

- is there a directed path from s to t ? (transitive closure)
- what is the shortest path from s to t ? (shortest path)
- are all vertices mutually reachable? (strong connectivity)
- how to organise a set of tasks? (topological sort)
- which web pages are "important"? (PageRank)
- how to build a web crawler? (graph traversal)


Digraph Representation

7/86

Similar set of choices as for undirectional graphs:

- array of edges (directed)
- vertex-indexed adjacency matrix (non-symmetric)
- vertex-indexed adjacency lists

V vertices identified by $0 \dots V-1$



Reachability

Transitive Closure

9/86

Given a digraph G it is potentially useful to know

- is vertex t reachable from vertex s ?

Example applications:

- can I complete a schedule from the current state?
- is a malloc'd object being referenced by any pointer?

How to compute transitive closure?

... Transitive Closure

10/86

One possibility:

- implement it via `hasPath(G, s, t)` (itself implemented by DFS or BFS algorithm)
- feasible if `reachable(G, s, t)` is infrequent operation

What if we have an algorithm that frequently needs to check reachability?

Would be very convenient/efficient to have:


```
reachable( $G, s, t$ ):
|   return  $G.tc[s][t]$     // transitive closure matrix
```

Of course, if V is very large, then this is not feasible.


Exercise #1: Transitive Closure Matrix

11/86

Which reachable $s \dots t$ exist in the following graph?



Transitive closure of example graph:



... Transitive Closure

13/86

Goal: produce a matrix of reachability values

- if $tc[s][t]$ is 1, then t is reachable from s
- if $tc[s][t]$ is 0, then t is not reachable from s

So, how to create this matrix?

Observation:

$\forall i, s, t \in \text{vertices}(G)$:

$(s, i) \in \text{edges}(G) \text{ and } (i, t) \in \text{edges}(G) \Rightarrow tc[s][t] = 1$

$tc[s][t] = 1$ if there is a path from s to t of length 2 ($s \rightarrow i \rightarrow t$)

... Transitive Closure

14/86

If we implement the above as:

```
make tc[][] a copy of edges[][]
for all i ∈ vertices(G) do
  for all s ∈ vertices(G) do
    for all t ∈ vertices(G) do
      if tc[s][i] = 1 and tc[i][t] = 1 then
        tc[s][t] = 1
      end if
    end for
  end for
end for
```

then we get an algorithm to convert `edges` into a `tc`

This is known as *Warshall's algorithm*

... Transitive Closure

15/86

How it works ...

After iteration 1, $tc[s][t]$ is 1 if

- either $s \rightarrow t$ exists or $s \rightarrow 0 \rightarrow t$ exists

After iteration 2, $tc[s][t]$ is 1 if any of the following exist

- $s \rightarrow t$ or $s \rightarrow 0 \rightarrow t$ or $s \rightarrow 1 \rightarrow t$ or $s \rightarrow 0 \rightarrow 1 \rightarrow t$ or $s \rightarrow 1 \rightarrow 0 \rightarrow t$


Etc. ... so after the V^{th} iteration, $tc[s][t]$ is 1 if

- there is any directed path in the graph from s to t

Exercise #2: Transitive Closure

16/86

Trace Warshall's algorithm on the following graph:



1st iteration $i=0$:

tc	[0]	[1]	[2]	[3]
[0]	0	1	1	1
[1]	1	1	1	1
[2]	0	1	0	0
[3]	0	0	0	0

2nd iteration $i=1$:

tc	[0]	[1]	[2]	[3]
[0]	1	1	1	1
[1]	1	1	1	1
[2]	1	1	1	1
[3]	0	0	0	0

3rd iteration $i=2$: unchanged

4th iteration $i=3$: unchanged

... Transitive Closure

18/86

Cost analysis:

- storage: additional V^2 items (each item may be 1 bit)
- computation of transitive closure: $O(V^3)$
- computation of `reachable()`: $O(I)$ after having generated `tc[][]`

Amortisation: would need many calls to `reachable()` to justify other costs

Alternative: use DFS in each call to `reachable()`

Cost analysis:

- storage: cost of queue and set during `reachable`
- computation of `reachable()`: cost of DFS = $O(V^2)$ (for adjacency matrix)

Digraph Traversal

19/86

Same algorithms as for undirected graphs:

depthFirst(v):

1. mark v as visited
2. for each $(v, w) \in \text{edges}(G)$ do
if w has not been visited then
depthFirst(w)

breadth-first(v):

1. enqueue v

2. while queue not empty do
 - dequeue v
 - if v not already visited then
 - mark v as visited
 - enqueue each vertex w adjacent to v

Example: Web Crawling

20/86

Goal: visit every page on the web

Solution: breadth-first search with "implicit" graph

```
webCrawl(startingURL):
|   mark startingURL as alreadySeen
|   enqueue(Q, startingURL)
|   while Q is not empty do
|       |   nextPage=dequeue(Q)
|       |   visit nextPage
|       |   for each hyperLink on nextPage do
|       |       |   if hyperLink not alreadySeen then
|       |       |       mark hyperLink as alreadySeen
|       |       |       enqueue(Q, hyperLink)
|       |       end if
|       end for
|   end while
```

visit scans page and collects e.g. keywords and links

Weighted Graphs

Weighted Graphs

22/86

Graphs so far have considered

- edge = an association between two vertices/nodes
- may be a precedence in the association (directed)

Some applications require us to consider


- a *cost* or *weight* of an association
- modelled by assigning values to edges (e.g. positive reals)

Weights can be used in both directed and undirected graphs.

... Weighted Graphs

23/86

Example: major airline flight routes in Australia



Representation: edge = direct flight; weight = approx flying time (hours)

... Weighted Graphs

24/86

Weights lead to minimisation-type questions, e.g.

1. Cheapest way to connect all vertices?

- a.k.a. *minimum spanning tree* problem
- assumes: edges are weighted and undirected

2. Cheapest way to get from *A* to *B*?

- a.k.a. *shortest path* problem
- assumes: edge weights positive, directed or undirected

Exercise #3: Implementing a Route Finder

25/86

If we represent a street map as a graph

- what are the vertices?
- what are the edges?
- are edges directional?
- what are the weights?
- are the weights fixed?

Weighted Graph Representation

26/86

Weights can easily be added to:


- adjacency matrix representation (0/1 \rightarrow int or float)
- adjacency lists representation (add int/float to list node)

Both representations work whether edges are directed or not.

... Weighted Graph Representation

27/86

Adjacency matrix representation with weights:




Note: need distinguished value to indicate "no edge".

... Weighted Graph Representation

28/86

Adjacency lists representation with weights:



Note: if undirected, each edge appears twice with same weight

... Weighted Graph Representation

29/86

Sample adjacency matrix implementation in C requires minimal changes to previous Graph ADT:

WGraph.h

```
// edges are pairs of vertices (end-points) plus positive weight
typedef struct Edge {
    Vertex v;
    Vertex w;
    int weight;
} Edge;

// returns weight, or 0 if vertices not adjacent
int adjacent(Graph, Vertex, Vertex);
```

... Weighted Graph Representation

30/86

WGraph.c

```
typedef struct GraphRep {
    int **edges; // adjacency matrix storing positive weights
                // 0 if nodes not adjacent
    int nV;      // #vertices
    int nE;      // #edges
} GraphRep;

void insertEdge(Graph g, Edge e) {
    assert(g != NULL && validV(g,e.v) && validV(g,e.w));
    if (g->edges[e.v][e.w] == 0) { // edge e not in graph
        g->edges[e.v][e.w] = e.weight;
        g->nE++;
    }
}
```



```

}

int adjacent(Graph g, Vertex v, Vertex w) {
    assert(g != NULL && validV(g,v) && validV(g,w));
    return g->edges[v][w];
}


```

Minimum Spanning Trees

Exercise #4: Minimising Wires in Circuits

32/86

Electronic circuit designs often need to make the pins of several components electrically equivalent by wiring them together.



To interconnect a set of n pins we can use an arrangement of $n-1$ wires each connecting two pins.

What kind of algorithm would ...

- help us find the arrangement with the least amount of wire?

Minimum Spanning Trees

33/86

Reminder: *Spanning tree* ST of graph $G=(V,E)$

- *spanning* = all vertices, *tree* = no cycles
 - ST is a subgraph of G ($G'=(V,E')$ where $E' \subseteq E$)
 - ST is *connected* and *acyclic* 不成cycle

Minimum spanning tree MST of graph G

- MST is a spanning tree of G
- sum of edge weights is no larger than any other ST

Applications: Computer networks, Electrical grids, Transportation networks ...


Problem: how to (efficiently) find MST for graph G ?

NB: MST may not be unique (e.g. all edges have same weight \Rightarrow every ST is MST)


... Minimum Spanning Trees

34/86

Example:



An MST ...



... Minimum Spanning Trees

35/86

Brute force solution:

```

findMST(G):
    Input  graph G
    Output a minimum spanning tree of G

    bestCost = ∞
    for all spanning trees t of G do
        if cost(t) < bestCost then
            bestTree = t
            bestCost = cost(t)
        end if
    end for
    return bestTree
  
```

Example of *generate-and-test* algorithm.

Not useful because #spanning trees is potentially large (e.g. n^{n-2} for a complete graph with n vertices)

... Minimum Spanning Trees

36/86

Simplifying assumption:

- edges in G are not directed (MST for digraphs is harder)

Kruskal's Algorithm

37/86

One approach to computing MST for graph G with V nodes:

1. start with empty MST
2. consider edges in increasing weight order

- add edge if it does not form a cycle in MST
3. repeat until $V-1$ edges are added


Critical operations:

- iterating over edges in weight order
- checking for cycles in a graph

... Kruskal's Algorithm

38/86


Execution trace of Kruskal's algorithm:




Exercise #5: Kruskal's Algorithm

39/86


Show how Kruskal's algorithm produces an MST on:




After 3rd iteration:




After 6th iteration:



After 7th iteration:



After 8th iteration ($V-1=8$ edges added):



... Kruskal's Algorithm

41/86

Pseudocode:

```

KruskalMST(G):
    Input  graph G with n nodes
    Output a minimum spanning tree of G

    MST = empty graph
    sort edges(G) by weight
    for each e in sortedEdgeList do
        MST = MST ∪ {e}
        if MST has a cycle then
            MST = MST \ {e}
        end if
        if MST has n-1 edges then
            return MST
        end if
    end for
  
```

... Kruskal's Algorithm

42/86

Rough time complexity analysis ...

- sorting edge list is $O(E \cdot \log E)$
- at least V iterations over sorted edges
- on each iteration ...
 - getting next lowest cost edge is $O(1)$
 - checking whether adding it forms a cycle: cost = ??

Possibilities for cycle checking:

- use DFS ... too expensive?
- could use *Union-Find data structure* (see Sedgewick Ch.1)

Prim's Algorithm

43/86

Another approach to computing MST for graph $G=(V,E)$:

1. start from any vertex v and empty MST

2. choose edge not already in MST to add to MST
 - must be incident on a vertex s already connected to v in MST
 - must be incident on a vertex t not already connected to v in MST
 - must have minimal weight of all such edges
3. repeat until MST covers all vertices


Critical operations:

- checking for vertex being connected in a graph
- finding min weight edge in a set of edges

... Prim's Algorithm

44/86


Execution trace of Prim's algorithm (starting at $s=0$):



Exercise #6: Prim's Algorithm


45/86

Show how Prim's algorithm produces an MST on:



Start from vertex 0


After 1st iteration:




After 2nd iteration:




After 3rd iteration:



After 4th iteration:



After 8th iteration (all vertices covered):



... Prim's Algorithm

47/86

Pseudocode:

```

PrimMST(G):
    Input  graph G with n nodes
    Output a minimum spanning tree of G

    MST = empty graph
    usedV = {0}
    unusedE = edges(g)
    while |usedV| < n do
        find e=(s,t,w) ∈ unusedE such that {
            s ∈ usedV, t ∉ usedV and w is min weight of all such edges
        }
        MST = MST ∪ {e}
        usedV = usedV ∪ {t}
        unusedE = unusedE \ {e}
    end while
    return MST
  
```

Critical operation: finding best edge

... Prim's Algorithm

48/86

Rough time complexity analysis ...

- V iterations of outer loop
- in each iteration ...
 - find min edge with set of edges is $O(E) \Rightarrow O(V \cdot E)$ overall
 - find min edge with *priority queue* is $O(\log E) \Rightarrow O(V \cdot \log E)$ overall

Sidetrack: Priority Queues

Some applications of queues require

- items processed in order of "priority"
- rather than in order of entry (FIFO — first in, first out)

Priority Queues (PQueues) provide this via:

- **join**: insert item into PQueue with an associated priority (replacing enqueue)
- **leave**: remove item with highest priority (replacing dequeue)

Time complexity for naive implementation of a PQueue containing N items ...

- $O(1)$ for join $O(N)$ for leave

Most efficient implementation ("heap") ...

- $O(\log N)$ for join, leave

Other MST Algorithms

50/86

Boruvka's algorithm ... complexity $O(E \cdot \log V)$

- the oldest MST algorithm
- start with V separate components
- join components using min cost links
- continue until only a single component

Karger, Klein, and Tarjan ... complexity $O(E)$

- based on Boruvka, but non-deterministic
- randomly selects subset of edges to consider
- for the keen, here's [the paper](#) describing the algorithm

Shortest Path

Shortest Path

52/86

Path = sequence of edges in graph G $p = (v_0, v_1), (v_1, v_2), \dots, (v_{m-1}, v_m)$

$cost(path)$ = sum of edge weights along path

Shortest path between vertices s and t

- a simple path $p(s, t)$ where $s = first(p)$, $t = last(p)$
- no other simple path $q(s, t)$ has $cost(q) < cost(p)$

Assumptions: weighted digraph, no negative weights.

Finding shortest path between two given nodes known as *source-target* SP problem

Variations: *single-source* SP, *all-pairs* SP

Applications: navigation, routing in data networks, ...

Single-source Shortest Path (SSSP)

53/86

Given: weighted digraph G , source vertex s

Result: shortest paths from s to all other vertices

- $\text{dist}[]$ V -indexed array of cost of shortest path from s
- $\text{pred}[]$ V -indexed array of predecessor in shortest path from s

Example:



Edge Relaxation


54/86

Assume: $\text{dist}[]$ and $\text{pred}[]$ as above (but containing data for shortest paths *discovered so far*)

$\text{dist}[v]$ is length of shortest known path from s to v

$\text{dist}[w]$ is length of shortest known path from s to w

Relaxation updates data for w if we find a shorter path from s to w :



Relaxation along edge $e=(v,w,\text{weight})$:

- if $\text{dist}[v] + \text{weight} < \text{dist}[w]$ then
update $\text{dist}[w] := \text{dist}[v] + \text{weight}$ and $\text{pred}[w] := v$

Dijkstra's Algorithm

55/86

One approach to solving single-source shortest path problem ...

Data: $G, s, \text{dist}[], \text{pred}[]$ and

- $vSet$: set of vertices whose shortest path from s is unknown

Algorithm:

```
dist[] // array of cost of shortest path from s
pred[] // array of predecessor in shortest path from s
```

```
dijkstraSSSP(G, source):
```


Input graph G , source node

initialise $\text{dist}[]$ to all ∞ , except $\text{dist}[\text{source}] = 0$

initialise $\text{pred}[]$ to all -1

$\text{vSet} = \text{all vertices of } G$

while $\text{vSet} \neq \emptyset$ **do**

find $s \in \text{vSet}$ **with** minimum $\text{dist}[s]$

for each $(s, t, w) \in \text{edges}(G)$ **do**

 relax along (s, t, w)

end for


$\text{vSet} = \text{vSet} \setminus \{s\}$

end while

Exercise #7: Dijkstra's Algorithm

56/86

Show how Dijkstra's algorithm runs on (source node = 0):



	[0]	[1]	[2]	[3]	[4]	[5]
dist	0	∞	∞	∞	∞	∞
pred	-	-	-	-	-	-

	[0]	[1]	[2]	[3]	[4]	[5]
dist	0	∞	∞	∞	∞	∞
pred	-	-	-	-	-	-

dist	0	14	9	7	∞	∞
pred	-	0	0	0	-	-

dist	0	14	9	7	∞	22
pred	-	0	0	0	-	3

dist	0	13	9	7	∞	12
pred	-	2	0	0	-	2

dist	0	13	9	7	20	12
pred	-	2	0	0	5	2

dist	0	13	9	7	18	12

pred	-	2	0	0	1	2
------	---	---	---	---	---	---

... Dijkstra's Algorithm

58/86

Why Dijkstra's algorithm is correct:

Hypothesis.

- (a) For visited s ... $dist[s]$ is shortest distance from source
- (b) For unvisited t ... $dist[t]$ is shortest distance from source *via visited nodes*

Proof.

Base case: no visited nodes, $dist[source]=0$, $dist[s]=\infty$ for all other nodes

Induction step:

1. If s is unvisited node with minimum $dist[s]$, then $dist[s]$ is shortest distance from source to s :
 - if \exists shorter path via only visited nodes, then $dist[s]$ would have been updated when processing the predecessor of s on this path
 - if \exists shorter path via an unvisited node u , then $dist[u] < dist[s]$, which is impossible if s has min distance of all unvisited nodes
2. This implies that (a) holds for s after processing s
3. (b) still holds for all unvisited nodes t after processing s :
 - if \exists shorter path via s we would have just updated $dist[t]$
 - if \exists shorter path without s we would have found it previously

... Dijkstra's Algorithm

59/86

Time complexity analysis ...

Each edge needs to be considered once $\Rightarrow O(E)$.

Outer loop has $O(V)$ iterations.

Implementing "**find** $s \in vSet$ **with** minimum $dist[s]$ "

1. try all $s \in vSet \Rightarrow cost = O(V) \Rightarrow overall\ cost = O(E + V^2) = O(V^2)$
2. using a PQueue to implement extracting minimum
 - can improve overall cost to $O(E + V \cdot \log V)$ (for best-known implementation)

All-pair Shortest Path (APSP)


60/86

Given: weighted digraph G

Result: shortest paths between all pairs of vertices

- $dist[][]$ $V \times V$ -indexed matrix of cost of shortest path from v_{row} to v_{col}
- $path[][]$ $V \times V$ -indexed matrix of next node in shortest path from v_{row} to v_{col}

Example:



V	0	1	2	3	4	
0	0	0.3	0.7	1.1	inf	dist
1	1.8	0	0.6	1.0	inf	
2	1.2	0.5	0	0.4	inf	
3	0.8	1.1	1.5	0	inf	
4	1.3	0.6	0.1	0.5	0	
0	-	1	2	2	-	path
1	2	-	2	2	-	
2	3	1	-	3	-	
3	0	0	0	-	-	
4	2	2	2	2	-	

Shortest paths between all vertices

Floyd's Algorithm

61/86

One approach to solving all-pair shortest path problem...

Data: G , $\text{dist}[][], \text{path}[][]$ Algorithm:

$\text{dist}[][]$ // array of cost of shortest path from s to t
 $\text{path}[][]$ // array of next node after s on shortest path from s to t

```

floydAPSP(G):
    Input graph G


    initialise dist[s][t]=0 for each s=t
                                =w for each (s,t,w)∈edges(G)
                                =∞ otherwise
    initialise path[s][t]=t for each (s,t,w)∈edges(G)
                                =-1 otherwise

    for all i∈vertices(G) do
        for all s∈vertices(G) do
            for all t∈vertices(G) do
                if dist[s][i]+dist[i][t] < dist[s][t] then
                    dist[s][t]=dist[s][i]+dist[i][t]
                    path[s][t]=path[s][i]
                end if
            end for
        end for
    end for
  
```

Exercise #8: Floyd's Algorithm

62/86

Show how Floyd's algorithm runs on:



dist	[0]	[1]	[2]	[3]	[4]	[5]	path	[0]	[1]	[2]	[3]	[4]	[5]
------	-----	-----	-----	-----	-----	-----	------	-----	-----	-----	-----	-----	-----

[0]	0	14	9	7			[0]		1	2	3		
[1]		0			5		[1]					4	
[2]		4	0			3	[2]		1				5
[3]			10	0		15	[3]			2			5
[4]					0		[4]						
[5]					2	0	[5]					4	

After 1st iteration i=0: unchanged After 2nd iteration i=1:

dist	[0]	[1]	[2]	[3]	[4]	[5]	path	[0]	[1]	[2]	[3]	[4]	[5]
[0]	0	14	9	7	19	∞	[0]	–	1	2	3	1	–
[1]	∞	0	∞	∞	5	∞	[1]	–	–	–	–	4	–
[2]	∞	4	0	∞	9	3	[2]	–	1	–	–	1	5
[3]	∞	∞	10	0	∞	15	[3]	–	–	2	–	–	5
[4]	∞	∞	∞	∞	0	∞	[4]	–	–	–	–	–	–
[5]	∞	∞	∞	∞	2	0	[5]	–	–	–	–	4	–

After 3rd iteration i=2:

dist	[0]	[1]	[2]	[3]	[4]	[5]	path	[0]	[1]	[2]	[3]	[4]	[5]
[0]	0	13	9	7	18	12	[0]	–	2	2	3	2	2
[1]	∞	0	∞	∞	5	∞	[1]	–	–	–	–	4	–
[2]	∞	4	0	∞	9	3	[2]	–	1	–	–	1	5
[3]	∞	14	10	0	19	13	[3]	–	2	2	–	2	2
[4]	∞	∞	∞	∞	0	∞	[4]	–	–	–	–	–	–
[5]	∞	∞	∞	∞	2	0	[5]	–	–	–	–	4	–

After 4th iteration i=3: unchanged After 5th iteration i=4: unchanged After 6th iteration i=5:


dist	[0]	[1]	[2]	[3]	[4]	[5]	path	[0]	[1]	[2]	[3]	[4]	[5]
[0]	0	13	9	7	14	12	[0]	–	2	2	3	2	2
[1]	∞	0	∞	∞	5	∞	[1]	–	–	–	–	4	–
[2]	∞	4	0	∞	5	3	[2]	–	1	–	–	5	5
[3]	∞	14	10	0	15	13	[3]	–	2	2	–	2	2
[4]	∞	∞	∞	∞	0	∞	[4]	–	–	–	–	–	–
[5]	∞	∞	∞	∞	2	0	[5]	–	–	–	–	4	–

... Floyd's Algorithm

Why Floyd's algorithm is correct:

A shortest path from s to t using only nodes from $\{0, \dots, i\}$ is the shorter of

- a shortest path from s to t using only nodes from $\{0, \dots, i-1\}$
- a shortest path from s to i using only nodes from $\{0, \dots, i-1\}$ plus a shortest path from i to t using only nodes from $\{0, \dots, i-1\}$



Also known as Floyd-Warshall algorithm (can you see why?)

... Floyd's Algorithm

65/86

Cost analysis ...

- initialising $\text{dist}[\][\], \text{path}[\][\] \Rightarrow O(E)$
- V iterations to update $\text{dist}[\][\], \text{path}[\][\] \Rightarrow O(V^3)$

Time complexity of Floyd's algorithm: $O(V^3)$ (same as Warshall's algorithm for transitive closure)


Network Flow

Exercise #9: Merchandise Distribution

67/86

Lucky Cricket Company ...

- produces cricket balls in Fairfield
- has a warehouse in Rozelle that stocks them
- ships them from factory to warehouse by leasing space on trucks with limited capacity:



What kind of algorithm would ...

- help us find the maximum number of crates that can be shipped from Fairfield to Rozelle per day?

Flow Networks

68/86

Flow network ...

- weighted graph $G=(V,E)$
- distinct nodes $s \in V$ (source), $t \in V$ (sink)

Edge weights denote *capacities* Applications:

- Distribution networks, e.g.
 - source: oil field
 - sink: refinery
 - edges: pipes
- Traffic flow

... Flow Networks

69/86

Flow in a network $G=(V,E)$... nonnegative $f(v,w)$ for all vertices $v,w \in V$ such that

- $f(v,w) \leq \text{capacity}$ for each edge $e=(v,w,\text{capacity}) \in E$
- $f(v,w)=0$ if no edge between v and w
- total flow *into* a vertex = total flow *out of* a vertex:


$$\sum_{x \in V} f(x, v) = \sum_{y \in V} f(v, y) \quad \text{for all } v \in V \setminus \{s, t\}$$

Maximum flow ... no other flow from s to t has larger value


... Flow Networks

70/86

Example:



A (maximum) flow ...




Augmenting Paths

71/86

Assume ... $f(v,w)$ contains current flow

Augmenting path: any path from source s to sink t that can currently take more flow

Example:



Residual Network


72/86

Assume ... flow network $G=(V,E)$ and flow $f(v,w)$

Residual network (V, E') :

- same vertex set V
- for each edge $v \xrightarrow{c} w \in E \dots$
 - $f(v, w) < c \Rightarrow$ add edge $(v \xrightarrow{c-f(v, w)} w)$ to E'
 - $f(v, w) > 0 \Rightarrow$ add edge $(v \xleftarrow{f(v, w)} w)$ to E'


Example:



Exercise #10: Augmenting Paths and Residual Networks


73/86

Find an augmenting path in:




and show the residual network after augmenting the flow

1. Augmenting path:



maximum additional flow = 1

2. Residual network:



Can you find a further augmenting path in the new residual network?

Edmonds-Karp Algorithm

75/86

One approach to solving maximum flow problem ...

maxflow(G):

1. Find a shortest augmenting path
2. Update flow[][] so as to represent residual graph
3. Repeat until no augmenting path can be found

... Edmonds-Karp Algorithm

76/86

Algorithm:

```
flow[ ][ ] // V×V array of current flow
visited[ ] /* array of predecessor nodes on shortest path
            from source to sink in residual network */
```

```
maxflow(G):
|   Input   flow network G with source s and sink t
|   Output maximum flow value
|
|   initialise flow[v][w]=0 for all vertices v, w
|   maxflow=0
|   while ∃ shortest augmenting path visited[ ] from s to t do
|       df = maximum additional flow via visited[ ]
|       // adjust flow so as to represent residual graph
|       v=t
|       while v≠s do
|           flow[visited[v]][v] = flow[visited[v]][v] + df;
|           flow[v][visited[v]] = flow[v][visited[v]] - df;
|           v=visited[v]
|       end while
|       maxflow=maxflow+df
|   end while
|   return maxflow
```

Shortest augmenting path can be found by standard BFS

... Edmonds-Karp Algorithm

77/86

Time complexity analysis ...

- *Theorem.* The number of augmenting paths needed is at most $V \cdot E/2$.
 \Rightarrow Outer loop has $O(V \cdot E)$ iterations.
- Finding augmenting path $\Rightarrow O(E)$ (consider only vertices connected to source and sink $\Rightarrow O(V+E)=O(E)$)


Overall cost of Edmonds-Karp algorithm: $O(V \cdot E^2)$

Note: Edmonds-Karp algorithm is an implementation of general *Ford-Fulkerson method*

Exercise #11: Edmonds-Karp Algorithm

78/86

Show how Edmonds-Karp algorithm runs on:



flow	[0]	[1]	[2]	[3]	[4]	[5]	c-f	[0]	[1]	[2]	[3]	[4]	[5]
[0]							[0]		2	3			
[1]							[1]				3	1	
[2]							[2]				1	1	
[3]							[3]						2
[4]							[4]						3
[5]							[5]						

flow	[0]	[1]	[2]	[3]	[4]	[5]	c-f	[0]	[1]	[2]	[3]	[4]	[5]
[0]	0	0	0	0	0	0	[0]	-	2	3	-	-	-
[1]	0	0	0	0	0	0	[1]	-	-	-	3	1	-
[2]	0	0	0	0	0	0	[2]	-	-	-	1	1	-
[3]	0	0	0	0	0	0	[3]	-	-	-	-	-	2
[4]	0	0	0	0	0	0	[4]	-	-	-	-	-	3
[5]	0	0	0	0	0	0	[5]	-	-	-	-	-	-

augmenting path: 0-1-3-5, df: 2

flow	[0]	[1]	[2]	[3]	[4]	[5]	c-f	[0]	[1]	[2]	[3]	[4]	[5]
[0]	0	2	0	0	0	0	[0]	-	0	3	-	-	-
[1]	-2	0	0	2	0	0	[1]	2	-	-	1	1	-
[2]	0	0	0	0	0	0	[2]	-	-	-	1	1	-
[3]	0	-2	0	0	0	2	[3]	-	2	-	-	-	0
[4]	0	0	0	0	0	0	[4]	-	-	-	-	-	3
[5]	0	0	0	-2	0	0	[5]	-	-	-	2	-	-

augmenting path: 0-2-4-5, df: 1

flow	[0]	[1]	[2]	[3]	[4]	[5]	c-f	[0]	[1]	[2]	[3]	[4]	[5]
[0]	0	2	1	0	0	0	[0]	-	0	2	-	-	-
[1]	-2	0	0	2	0	0	[1]	2	-	-	1	1	-
[2]	-1	0	0	0	1	0	[2]	1	-	-	1	0	-

[3]	0	-2	0	0	0	2	[3]	-	2	-	-	-	0
[4]	0	0	-1	0	0	1	[4]	-	-	1	-	-	2
[5]	0	0	0	-2	-1	0	[5]	-	-	-	2	1	-

augmenting path: 0-2-3-1-4-5, df: 1

flow	[0]	[1]	[2]	[3]	[4]	[5]	c-f	[0]	[1]	[2]	[3]	[4]	[5]
[0]	0	2	2	0	0	0	[0]	-	0	1	-	-	-
[1]	-2	0	0	1	1	0	[1]	2	-	-	2	0	-
[2]	-2	0	0	1	1	0	[2]	2	-	-	0	0	-
[3]	0	-1	-1	0	0	2	[3]	-	1	1	-	-	0
[4]	0	-1	-1	0	0	2	[4]	-	1	1	-	-	1
[5]	0	0	0	-2	-2	0	[5]	-	-	-	2	2	-

Digraph Applications

PageRank

81/86

Goal: determine which Web pages are "important"

Approach: ignore page contents; focus on hyperlinks

- treat Web as graph: page = vertex, hyperlink = directed edge
- pages with many incoming hyperlinks are important
- need to compute "incoming degree" for vertices

Problem: the Web is a *very* large graph

- approx. 10^{14} pages, 10^{15} hyperlinks

Assume for the moment that we could build a graph ...

Most frequent operation in algorithm "Does edge (v,w) exist?"

... PageRank

82/86

Simple PageRank algorithm:

```
PageRank(myPage):
|   rank=0
|   for each page in the Web do
|       if linkExists(page,myPage) then
|           rank=rank+1
|       end if
|   end for
```

Note: requires *inbound* link check

... PageRank

$V = \# \text{ pages in Web}$, $E = \# \text{ hyperlinks in Web}$

Costs for computing PageRank for each representation:

Representation	linkExists(v,w)	Cost
Adjacency <i>matrix</i>	edge[v][w]	I
Adjacency <i>lists</i>	inLL(list[v],w)	$\cong E/V$

Not feasible ...

- adjacency matrix ... $V \cong 10^{14} \Rightarrow$ matrix has 10^{28} cells
- adjacency list ... V lists, each with $\cong 10$ hyperlinks $\Rightarrow 10^{15}$ list nodes

So how to really do it?

... PageRank

Approach: the random web surfer

- if we randomly follow links in the web ...
- ... more likely to re-discover pages with many inbound links

```
curr=random page, prev=null
for a long time do
  if curr not in array ranked[] then
    rank[curr]=0
  end if
  rank[curr]=rank[curr]+1
  if random(0,100)<85 then           // with 85% chance ...
    prev=curr
    curr=choose hyperlink from curr // ... crawl on
  else
    curr=random page                // avoid getting stuck
    prev=null
  end if
end for
```

Could be accomplished while we crawl web to build search index

Exercise #12: Implementing Facebook

Facebook could be considered as a giant "social graph"

- what are the vertices?
- what are the edges?
- are edges directional?

What kind of algorithm would ...

- help us find people that you might like to "befriend"?

Summary

- Digraphs, weighted graphs: representations, applications
- Reachability
 - Warshall
- Minimum Spanning Tree (MST)
 - Kruskal, Prim
- Shortest path problems
 - Dijkstra (single source SPP)
 - Floyd (all-pair SSP)
- Flow networks
 - Edmonds-Karp (maximum flow)

- Suggested reading (Sedgewick):
 - digraphs ... Ch. 19.1-19.3
 - weighted graphs ... Ch. 20-20.1
 - MST ... Ch. 20.2-20.4
 - SSP ... Ch. 21-21.3
 - network flows ... Ch. 22.1-22.2

Produced: 26 Jun 2020