

## Week 8: Search Tree Algorithms

### Tree Review

1/72

Binary search trees ...

- data structures designed for  $O(\log n)$  search
- consist of nodes containing item (incl. key) and two links
- can be viewed as recursive data structure (subtrees)
- have overall ordering ( $\text{data}(\text{Left}) < \text{root} < \text{data}(\text{Right})$ )
- insert new nodes as leaves (or as root), delete from anywhere
- have structure determined by insertion order (*worst*:  $O(n)$ )
- operations: insert, delete, search, ...


### Balanced Search Trees

### Balanced BSTs

3/72

Reminder ...

- Goal: build binary search trees which have
  - minimum height  $\Rightarrow$  minimum worst case search cost
- Best balance you can achieve for tree with  $N$  nodes:
  - tree height of  $\log_2 N \Rightarrow$  worst case search  $O(\log N)$



Three *strategies* to improving worst case search in BSTs:

- *randomise* — reduce chance of worst-case scenario occurring
- *amortise* — do more work at insertion to make search faster
- *optimise* — implement all operations with performance bounds

### Randomised BST Insertion

4/72

Effects of order of insertion on BST shape:

- best case (for at-leaf insertion): keys inserted in pre-order  
(median key first, then median of lower half, median of upper half, etc.)

- worst case: keys inserted in ascending/descending order
- average case: keys inserted in *random* order  $\Rightarrow O(\log_2 n)$

Tree ADT has no control over order that keys are supplied.

Can the algorithm itself introduce some *randomness*?

In the hope that this randomness helps to balance the tree ...

## ... Randomised BST Insertion

5/72

How can a computer pick a number at random?

- it cannot

Software can only produce *pseudo random numbers*.

- a pseudo random number may appear unpredictable
  - but is actually predictable
- $\Rightarrow$  implementation may deviate from expected theoretical behaviour
  - more on this in week 10

## ... Randomised BST Insertion

6/72

- Pseudo random numbers in C:

```
rand() // generates random numbers in the range 0 .. RAND_MAX
```

where the constant `RAND_MAX` is defined in `stdlib.h`

(depends on the computer: on the CSE network, `RAND_MAX` = 2147483647)

To convert the return value of `rand()` to a number between 0 .. RANGE

- compute the remainder after division by `RANGE+1`

## ... Randomised BST Insertion

7/72

Approach: normally do leaf insert, randomly do root insert.

```
insertRandom(tree,item)
| Input tree, item
| Output tree with item randomly inserted
|
| if tree is empty then
|   return new node containing item
| end if
| // p/q chance of doing root insert
| if random number mod q < p then
|   return insertAtRoot(tree,item)
| else
|   return insertAtLeaf(tree,item)
| end if
```

E.g. 30% chance  $\Rightarrow$  choose  $p=3, q=10$

## ... Randomised BST Insertion

Cost analysis:

- similar to cost for inserting keys in random order:  $O(\log_2 n)$
- does not rely on keys being supplied in random order

Approach can also be applied to deletion:

- standard method promotes inorder successor to root
- for the randomised method ...
  - promote inorder successor from right subtree, OR
  - promote inorder predecessor from left subtree

## Rebalancing Trees

An approach to balanced trees:

- insert into leaves as for simple BST
- periodically, rebalance the tree

Question: how frequently/when/how to rebalance?

NewTreeInsert(tree, item):

```

Input   tree, item
Output tree with item randomly inserted

t=insertAtLeaf(tree,item)
if #nodes(t) mod k = 0 then
    t=rebalance(t)
end if
return t
  
```

E.g. rebalance after every 20 insertions  $\Rightarrow$  choose  $k=20$


Note: To do this efficiently we would need to change tree data structure and basic operations:

```

typedef struct Node {
    int data;
    int nnodes; // #nodes in my tree
    Tree left, right; // subtrees
} Node;
  
```

## ... Rebalancing Trees

How to rebalance a BST? Move median item to root.



## ... Rebalancing Trees

11/72

Implementation of rebalance:

**rebalance(t):**

**Input** tree t with n nodes

**Output** t rebalanced

**if**  $n \geq 3$  **then**

    t = **partition**(t,  $\lfloor n/2 \rfloor$ )      // put node with median key at root

    left(t) = **rebalance**(left(t))    // then rebalance each subtree

    right(t) = **rebalance**(right(t))

**end if**


return t

## ... Rebalancing Trees

12/72

New operation on trees:

- **partition(tree, i)**: re-arrange tree so that element with index  $i$  becomes root




For tree with  $N$  nodes, indices are  $0 \dots N-1$

## ... Rebalancing Trees

13/72

Partition: moves  $i^{\text{th}}$  node to root



## ... Rebalancing Trees

14/72

Implementation of partition operation:

**partition(tree, i):**

**Input** tree with n nodes, index  $i$

**Output** tree with item # $i$  moved to the root

```

m=#nodes(left(tree))
if i < m then
    left(tree)=partition(left(tree),i)
    tree=rotateRight(tree)
else if i > m then
    right(tree)=partition(right(tree),i-m-1)
    tree=rotateLeft(tree)
end if
return tree


```

Note:  $\text{size}(\text{tree}) = n$ ,  $\text{size}(\text{left}(\text{tree})) = m$ ,  $\text{size}(\text{right}(\text{tree})) = n-m-1$  (why -1?)


## Exercise #1: Partition

15/72

Consider the tree  $t$ :



Show the result of  $\text{partition}(t, 3)$



## ... Rebalancing Trees

17/72

Analysis of rebalancing: visits every node  $\Rightarrow O(N)$

Cost means not feasible to rebalance after each insertion.

When to rebalance? ... Some possibilities:

- after every  $k$  insertions
- whenever "imbalance" exceeds threshold

Either way, we tolerate worse search performance for periods of time.

Does it solve the problem? ... Not completely  $\Rightarrow$  Solution: real balanced trees (later)

## Splay Trees

## Splay Trees

A kind of "self-balancing" tree ...

Splay tree insertion modifies insertion-at-root method:

- by considering *parent-child-grandchild* (three level analysis)
- by performing double-rotations based on p-c-g orientation

The idea: appropriate double-rotations improve tree balance.

### ... Splay Trees

20/72

Splay tree implementations also do *rotation-in-search*:

- by performing double-rotations also when searching

The idea: provides similar effect to periodic rebalance.


⇒ improves balance but makes search more expensive

### ... Splay Trees

21/72

Cases for splay tree double-rotations:


- case 1: grandchild is left-child of left-child ⇒ double right rotation from top
- case 2: grandchild is right-child of left-child
- case 3: grandchild is left-child of right-child
- case 4: grandchild is right-child of right-child ⇒ double left rotation from top



### ... Splay Trees

22/72

Double-rotation case for left-child of left-child ("zig-zig"):




Note: both rotations at the root (unlike insertion-at-root)

## ... Splay Trees

23/72

Double-rotation case for right-child of left-child ("zig-zag"):



Note: rotate subtree first (like insertion-at-root)

## ... Splay Trees

24/72

Algorithm for splay tree insertion:

**insertSplay(tree, item):**

**Input** tree, item

**Output** tree with item splay-inserted

**if** tree is empty **then return** new node containing item

**else if** item=data(tree) **then return** tree

**else if** item<data(tree) **then**

**if** left(tree) is empty **then**

        left(tree)=new node containing item

**else if** item<data(left(tree)) **then**

        // Case 1: left-child of left-child "zig-zig"

        left(left(tree))=insertSplay(left(left(tree)),item)

        tree=rotateRight(tree)

**else if** item>data(left(tree)) **then**

        // Case 2: right-child of left-child "zig-zag"

        right(left(tree))=insertSplay(right(left(tree)),item)

        left(tree)=rotateLeft(left(tree))

**end if**

**return** rotateRight(tree)

**else** // item>data(tree)

**if** right(tree) is empty **then**

        right(tree)=new node containing item

**else if** item<data(right(tree)) **then**

        // Case 3: left-child of right-child "zag-zig"

        left(right(tree))=insertSplay(left(right(tree)),item)

        right(tree)=rotateRight(right(tree))

**else if** item>data(right(tree)) **then**

        // Case 4: right-child of right-child "zag-zag"

        right(right(tree))=insertSplay(right(right(tree)),item)

        tree=rotateLeft(tree)



**end if**

**return** rotateLeft(tree)

**end if**

## Exercise #2: Splay Trees

Insert 18 into this splay tree:



## ... Splay Trees

Searching in splay trees:

```


searchSplay(tree,item):
  Input  tree, item
  Output address of item if found in tree
          NULL otherwise

  if tree=NULL then
    return NULL
  else
    tree=splay(tree,item)
    if data(tree)=item then
      return tree
    else
      return NULL
    end if
  end if
  
```

where `splay()` is similar to `insertSplay()`, except that it doesn't add a node ... simply moves `item` to root if found, or nearest node if not found


## Exercise #3: Splay Trees

If we search for 22 in the splay tree





... how does this affect the tree?



## ... Splay Trees

30/72

Why take into account both child and grandchild?

- moves accessed node to the root
- *moves every ancestor of accessed node roughly halfway to the root*

⇒ better amortized cost than insert-at-root

## ... Splay Trees

31/72

Analysis of splay tree performance:

- assume that we "splay" for both insert and search
- consider:  $m$  insert+search operations,  $n$  nodes
- *Theorem.* Total number of comparisons: average  $O((n+m) \cdot \log(n+m))$

Gives good overall (amortized) cost.

- insert cost not significantly different to insert-at-root
- search cost increases, but ...
  - improves balance on each search
  - moves frequently accessed nodes closer to root

But ... still has worst-case search cost  $O(n)$

## Real Balanced Trees

### Better Balanced Binary Search Trees

33/72

So far, we have seen ...

- randomised trees ... make poor performance unlikely
- occasional rebalance ... fix balance periodically
- splay trees ... reasonable amortized performance
- but both types still have  $O(n)$  worst case

Ideally, we want both average/worst case to be  $O(\log n)$

- AVL trees ... fix imbalances as soon as they occur
- 2-3-4 trees ... use varying-sized nodes to assist balance
- red-black trees ... isomorphic to 2-3-4, but binary nodes

# AVL Trees

## AVL Trees

35/72

Invented by Georgy Adelson-Velsky and Evgenii Landis

Approach:

- insertion (at leaves) may cause imbalance
- repair balance as soon as we notice imbalance
- repairs done locally, not by overall tree restructure

A tree is unbalanced when:  $\text{abs}(\text{height}(\text{left}) - \text{height}(\text{right})) > 1$

This can be repaired by at most two rotations:

- if left subtree too deep ...
  - if data inserted in left-right grandchild  $\Rightarrow$  left-rotate left subtree
  - rotate right
- if right subtree too deep ...
  - if data inserted in right-left grandchild  $\Rightarrow$  right-rotate right subtree
  - rotate left

Problem: determining height/depth of subtrees may be expensive.

## ... AVL Trees

36/72



Implementation of AVL insertion

```
insertAVL(tree,item):
  Input  tree, item
  Output tree with item AVL-inserted

  if tree is empty then
    return new node containing item
  else if item=data(tree) then
    return tree
  else
    if item<data(tree) then
      left(tree)=insertAVL(left(tree),item)
    else if item>data(tree) then
      right(tree)=insertAVL(right(tree),item)
    end if
    if height(left(tree))-height(right(tree)) > 1 then
      if item>data(left(tree)) then
        left(tree)=rotateLeft(left(tree))
      end if
      tree=rotateRight(tree)
    else if height(right(tree))-height(left(tree)) > 1 then
      if item<data(right(tree)) then
        right(tree)=rotateRight(right(tree))
      end if
      tree=rotateLeft(tree)
    end if
    return tree
  end if
```

## Exercise #4: AVL Trees

Insert 27 into the AVL tree




What would happen if you now insert 28?

You may like the animation at [www.cs.usfca.edu/~galles/visualization/AVLtree.html](http://www.cs.usfca.edu/~galles/visualization/AVLtree.html)

## ... AVL Trees

Analysis of AVL trees:

- trees are *height*-balanced; subtree depths differ by  $\pm 1$
- average/worst-case search performance of  $O(\log n)$
- *require* extra data to be stored in each node ("height")
- may not be *weight*-balanced; subtree sizes may differ




## 2-3-4 Trees

## 2-3-4 Trees

2-3-4 trees have three kinds of nodes

- 2-nodes, with two children (same as normal BSTs)
- 3-nodes, two values and three children
- 4-nodes, three values and four children



## ... 2-3-4 Trees

42/72

2-3-4 trees are ordered similarly to BSTs



In a *balanced* 2-3-4 tree:

- all leaves are at same distance from the root

2-3-4 trees grow "upwards" by splitting 4-nodes.

## ... 2-3-4 Trees

43/72

Possible 2-3-4 tree data structure:

```

typedef struct node {
    int      order;      // 2, 3 or 4
    int      data[3];    // items in node
    struct node *child[4]; // links to subtrees
} node;
  
```

## ... 2-3-4 Trees

44/72

Searching in 2-3-4 trees:

**Search(tree,item):**

```

Input tree, item
Output address of item if found in 2-3-4 tree
        NULL otherwise

if tree is empty then
    return NULL
else
    i=0
    while i<tree.order-1 and item>tree.data[i] do
        i=i+1 // find relevant slot in data[]
    end while
    if item=tree.data[i] then // item found
        return address of tree.data[i]
    else // keep looking in relevant subtree
        return Search(tree.child[i],item)
    end if
end if
  
```

45/72

## ... 2-3-4 Trees

2-3-4 tree searching cost analysis:

- as for other trees, worst case determined by height  $h$
- 2-3-4 trees are always balanced  $\Rightarrow$  height is  $O(\log n)$
- worst case for height: all nodes are 2-nodes  
same case as for balanced BSTs, i.e.  $h \cong \log_2 n$
- best case for height: all nodes are 4-nodes  
balanced tree with branching factor 4, i.e.  $h \cong \log_4 n$

## Insertion into 2-3-4 Trees

46/72

Starting with the root node:

**repeat**


- if current node is full (i.e. contains 3 items)
  - split into two 2-nodes
  - promote middle element to parent
    - if no parent  $\Rightarrow$  middle element becomes the new root 2-node
  - go back to parent node
- if current node is a leaf
  - insert Item in this node, order++
- if current node is not a leaf
  - go to child where Item belongs

**until** Item inserted

## ... Insertion into 2-3-4 Trees

47/72

Insertion into a 2-node or 3-node:




Insertion into a 4-node (requires a split):



## ... Insertion into 2-3-4 Trees

48/72


Splitting the root:



## ... Insertion into 2-3-4 Trees

49/72



Building a 2-3-4 tree ... 7 insertions:



## Exercise #5: Insertion into 2-3-4 Tree

50/72

Show what happens when D, S, F, U are inserted into this tree:



## ... Insertion into 2-3-4 Trees

52/72

Insertion algorithm:

```

insert(tree, item):
    Input  2-3-4 tree, item
    Output tree with item inserted

    node=root(tree), parent=NULL
    repeat
        if node.order=4 then
            promote = node.data[1]      // middle value
            nodeL   = new node containing node.data[0]
            nodeR   = new node containing node.data[2]
            if parent=NULL then
                make new 2-node root with promote, nodeL, nodeR

```

```

else
    insert promote,nodeL,nodeR into parent
    increment parent.order
end if
node=parent
end if
if node is a leaf then
    insert item into node
    increment node.order
else
    parent=node
    if item<node.data[0] then
        node=node.child[0]
    else if item<node.data[1] then
        node=node.child[1]
    else
        node=node.child[2]
    end if
end if
until item inserted

```

## ... Insertion into 2-3-4 Trees

53/72

Variations on 2-3-4 trees ...

Variation #1: why stop at 4? why not 2-3-4-5 trees? or  $M$ -way trees?

- allow nodes to hold up to  $M-1$  items, and at least  $M/2$
- if each node is a disk-page, then we have a *B-tree* (databases)
- for B-trees, depending on Item size,  $M > 100/200/400$

Variation #2: don't have "variable-sized" nodes

- use standard BST nodes, augmented with one extra piece of data
- implement similar strategy as 2-3-4 trees → red-black trees.

## Red-Black Trees

## Red-Black Trees

55/72

*Red-black trees* are a representation of 2-3-4 trees using BST nodes.

- each node needs one extra value to encode link type
- but we no longer have to deal with different kinds of nodes

Link types:

- *red* links ... combine nodes to represent 3- and 4-nodes
- *black* links ... analogous to "ordinary" BST links (child links)

Advantages:

- standard BST search procedure works unmodified
- get benefits of 2-3-4 tree self-balancing (although deeper)

## Red-Black Trees

56/72

## Definition of a *red-black tree*

- a BST in which each node is marked red or black
- no two red nodes appear consecutively on any path
- a red node corresponds to a 2-3-4 sibling of its parent
- a black node corresponds to a 2-3-4 child of its parent

## Balanced red-black tree

- all paths from root to leaf have same number of black nodes


Insertion algorithm: avoids worst case  $O(n)$  behaviour

Search algorithm: standard BST search

## ... Red-Black Trees

57/72

Representing 4-nodes in red-black trees:




Some texts colour the links rather than the nodes.

## ... Red-Black Trees

58/72

Representing 3-nodes in red-black trees (two possibilities):



## ... Red-Black Trees

59/72

Equivalent trees (one 2-3-4, one red-black):





### ... Red-Black Trees

60/72

Red-black tree implementation:

```
typedef enum {RED, BLACK} Colr;
typedef struct node *RBTree;
typedef struct node {
    int    data;    // actual data
    Colr   colour;  // relationship to parent
    RBTree left;    // left subtree
    RBTree right;   // right subtree
} node;

#define colour(tree) ((tree)->colour)
#define isRed(tree)  ((tree) != NULL && (tree)->colour == RED)
```

RED = node is part of the same 2-3-4 node as its parent (sibling)

BLACK = node is a child of the 2-3-4 node containing the parent

### ... Red-Black Trees

61/72

New nodes are always red:


```
RBTree newNode(Item it) {
    RBTree new = malloc(sizeof(Node));
    assert(new != NULL);
    data(new) = it;
    colour(new) = RED;
    left(new) = right(new) = NULL;
    return new;
}
```

### ... Red-Black Trees

62/72

Node.colour allows us to distinguish links

- black = parent node is a "real" parent
- red = parent node is a 2-3-4 neighbour



## ... Red-Black Trees

63/72

Search method is standard BST search:

**SearchRedBlack(tree, item):**

```

Input  tree, item
Output true if item found in red-black tree
         false otherwise

if tree is empty then
    return false
else if item < data(tree) then
    return SearchRedBlack(left(tree), item)
else if item > data(tree) then
    return SearchRedBlack(right(tree), item)
else // found
    return true
end if

```

## Red-Black Tree Insertion

64/72

Insertion is more complex than for standard BSTs

- need to recall direction of last branch (L or R)
- need to recall whether parent link is red or black
- splitting/promoting implemented by rotateLeft/rotateRight
- several cases to consider depending on colour/direction combinations

## ... Red-Black Tree Insertion

65/72

High-level description of insertion algorithm:

**insertRB(tree, item, inRight):**

```

Input  tree, item, inRight indicating direction of last branch
Output tree with it inserted

if tree is empty then
    return newNode(item)
else if item = data(tree) then
    return tree
end if
if left(tree) and right(tree) both are RED then
    split 4-node in a red-black tree
end if
recursive insert a la BST, re-arrange links/colours after insert
return modified tree

```

```

insertRedBlack(tree,item):
|   Input   red-black tree, item
|   Output tree with item inserted
|
|   tree=insertRB(tree,item,false)
|   colour(tree)=BLACK
|   return tree


```

---

### ... Red-Black Tree Insertion

66/72

Splitting a 4-node, in a red-black tree:



Algorithm:

```

colour(currentTree)=RED
colour(left(currentTree))=BLACK
colour(right(currentTree))=BLACK

```

---

### ... Red-Black Tree Insertion

67/72

Simple recursive insert (a la BST):



Algorithm:

```

if item<data(tree) then
|   left(tree)=insertRB(left(tree),item,false)
|   re-arrange links/colours after insert
else           // item larger than data in root
|   right(tree)=insertRB(right(tree),item,true)
|   re-arrange links/colours after insert
end if

```

Not affected by colour of tree node.

---

### ... Red-Black Tree Insertion

68/72

Re-arrange links/colours after insert:

Step 1 — "normalise" direction of successive red links



Algorithm:

```

if inRight and both currentTree and left(currentTree) are red then
    currentTree=rotateRight(currentTree)
end if

```

Symmetrically,


- if not inRight and both currentTree and right(currentTree) are red  
⇒ left rotate currentTree

## ... Red-Black Tree Insertion

69/72

Re-arrange links/colours after insert:

Step 2 — two successive red links = newly-created 4-node



Algorithm:

```

if left(currentTree) and left(left(currentTree)) are red then
    currentTree=rotateRight(currentTree)
    colour(currentTree)=BLACK
    colour(right(currentTree))=RED
end if

```

Symmetrically,

- if both right(currentTree) and right(right(currentTree)) are red  
⇒ left rotate currentTree, then re-colour currentTree and left(currentTree)

## ... Red-Black Tree Insertion

70/72

Example of insertion, starting from empty tree:

22, 12, 8, 15, 11, 19, 43, 44, 45, 42, 41, 40, 39



## Red-black Tree Performance

71/72

Cost analysis for red-black trees:

- tree is well-balanced; worst case search is  $O(\log_2 n)$
- insertion affects nodes down one path; max #rotations/recolourings is  $O(h)$   
(where  $h$  is the height of the tree)

Only disadvantage is complexity of insertion/deletion code.

Note: red-black trees were popularised by Sedgewick.

---

## Summary

72/72

- Randomised at-leaf/at-root insertion
  - Tree operations
    - tree partition
    - joining trees
  - Self-adjusting trees
    - Splay trees
    - AVL trees
    - 2-3-4 trees
    - Red-black trees
  - Suggested reading:
    - Sedgewick, Ch. 12.9
    - Sedgewick, Ch. 13.1-13.4
- 

Produced: 21 Jul 2020